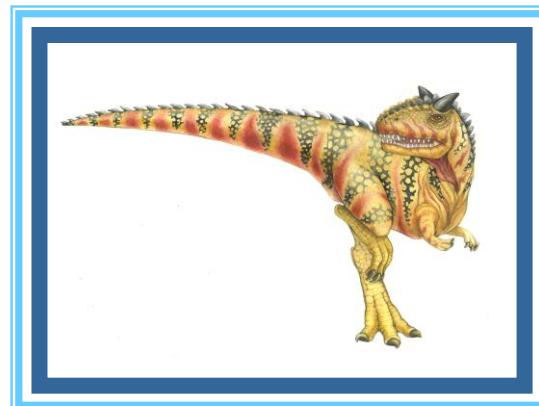


Chapter 6: Synchronization Tools





Chapter 6: Synchronization Tools

Background

The Critical-Section Problem

Peterson's Solution

Hardware Support for Synchronization

Mutex Locks

Semaphores

Monitors

Liveness

Evaluation





Objectives

Describe the **critical-section problem** and illustrate a **race condition**

Illustrate hardware solutions to the critical-section problem using memory barriers, compare-and-swap operations, and atomic variables

Demonstrate how **mutex locks**, **semaphores**, **monitors**, and **condition variables** can be used to solve the critical section problem

Evaluate tools that solve the critical-section problem in low-. Moderate-, and high-contention scenarios





Background

Processes can execute concurrently

May be interrupted at any time, partially completing execution

Concurrent access to shared data may result in data inconsistency

Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

prevent concurrent execution

Illustration of the problem:

Shared variable

Suppose that we wanted to provide a solution to the consumer-producer problem that fills all the buffers. We can do so by having an integer counter that keeps track of the number of full buffers. Initially, counter is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





Producer

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER_SIZE) ⇒ full buffer  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```





* Producer

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
counter++;  
}
```

* Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
counter--;  
    /* consume the item in next consumed */  
}
```

possible to access the variable
concurrently by different processes





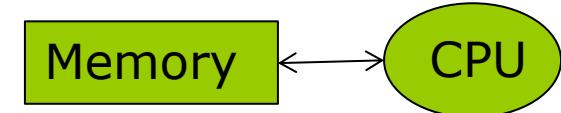
Race Condition

producer `counter++` could be implemented as

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

consumer counter-- could be implemented as

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```



depending on the order of execution of statement,
the result of executing two statement is unstable.

In race condition, the result of execution is unpredictable.
⇒ difficult to implement function

Consider this execution interleaving with “count = 5” initially:

```
S0: producer execute register1 = counter      {register1 = 5}
S1: producer execute register1 = register1 + 1 {register1 = 6}
S2: consumer execute register2 = counter       {register2 = 5}
S3: consumer execute register2 = register2 - 1 {register2 = 4}
S4: producer execute counter = register1        {counter = 6 }
S5: consumer execute counter = register2        {counter = 4}
                                         error !
```

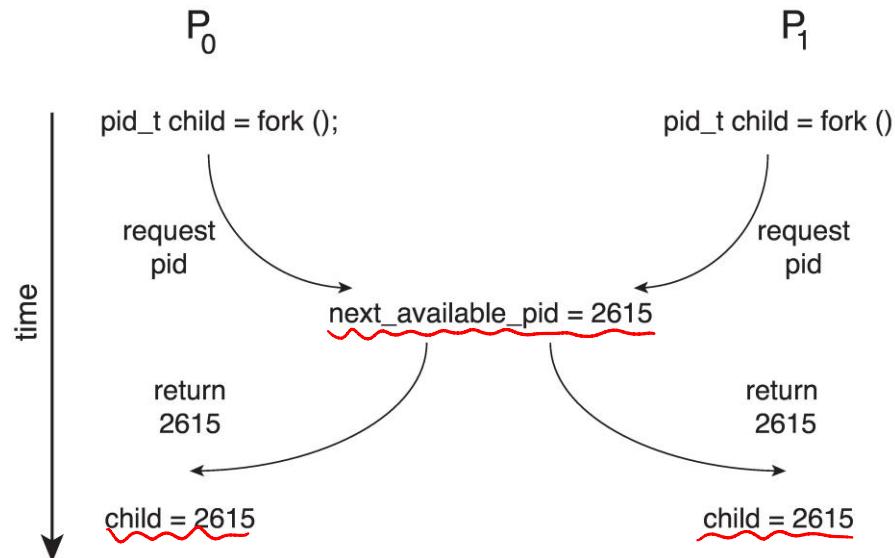




Race Condition (Cont'd)

Processes P_0 and P_1 are creating child processes using the `fork()` system call

Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)



Unless there is mutual exclusion, the same pid could be assigned to two different processes!

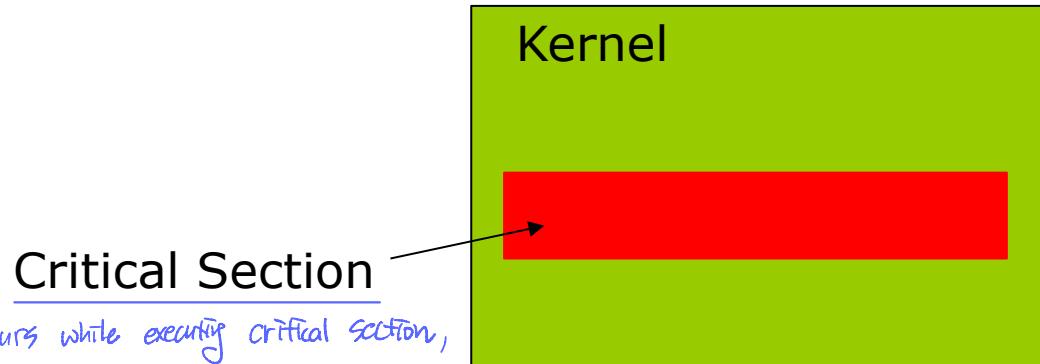




Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- Preemptive** – allows preemption of process when running in kernel mode *interrupt / context switching is allowed \Rightarrow race condition!* *handling critical-section problem*
- Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU *no interrupt, context switching \Rightarrow free from race condition*
- ▶ Essentially free of race conditions in kernel mode



*if preemption occurs while executing critical section,
race condition can occur
 \Rightarrow guarantees correct execution by mutual exclusion
for preservation of data consistency*





Critical Section Problem

Consider system of n processes $\{p_0, p_1, \dots p_{n-1}\}$

Each process has **critical section** segment of code

Process may be changing common variables, updating table, writing file, etc

When one process in critical section, no other may be in its critical section *mutual exclusion*

Critical section problem is to design protocol to solve this

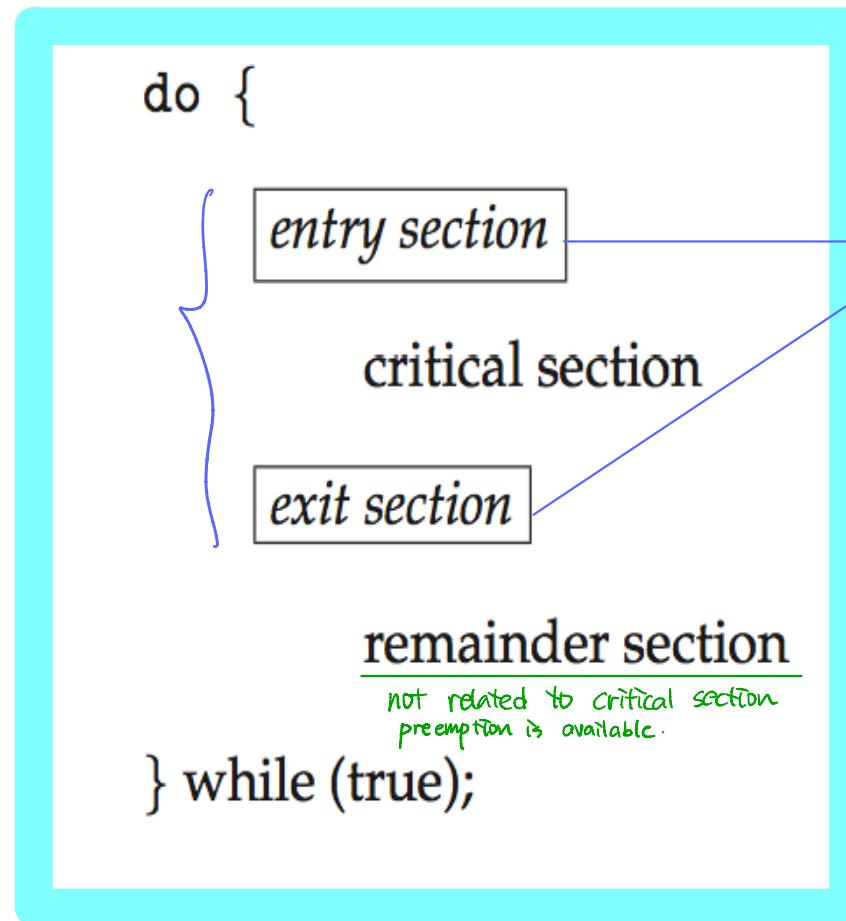
Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





Critical Section

General structure of process P_i



guarantee no other process
can access critical section
while the process is executing critical
section





Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
no preemption now, critical section must be guaranteed immediately
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning relative speed of the n processes
only limited number of processes can enter before the process gets its turn



Peterson's Solution

Not guaranteed to work on modern architectures! (But good algorithmic description of solving the problem)

Two process solution $P_i \& P_j$

Assume that the **load** and **store** machine-language instructions are **atomic**; that is, cannot be interrupted

The two processes share two variables:

```
int turn;  
boolean flag[2]
```

The variable turn indicates whose turn it is to enter the critical section
 $= i \text{ or } j$

The flag array is used to indicate if a process is ready to enter the critical section. flag[i] = true implies that process P_i is ready!





Algorithm for Process P_i (P_j)

P_i

```
while (true) {
    entry section } flag[i] = true;
    turn = j; yield chance to j if j want
    while (flag[j] && turn == j)
        ;
    /* critical section */
    exit section } flag[i] = false;
    /* remainder section */
}
```

P_j

```
while (true) {
    flag[j] = true;
    turn = i; Cannot be true at the same time
    while (flag[i] && turn == i)
        ;
    /* critical section */
    flag[j] = false;
    /* remainder section */
}
```

- * What happens if flag[] is not used?? (1)
- * What happens if turn is not used?? (2)





Algorithm for Process P_i (P_j)

P_i

```
while (true) {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;
    /* critical section */

    flag[i] = false;

    /* remainder section */
}
```

P_j

```
while (true) {
    flag[j] = true;
    turn = i;
    while (flag[i] && turn == i)
        ;
    /* critical section */

    flag[j] = false;  $\Rightarrow$  bounded waiting

    /* remainder section */
}

who executed first is depending on turn = i / turn = j execute first.

after yielding, progress is guaranteed.

progress is guaranteed directly
```

- * What happens if flag[] is not used?? (1)
- * What happens if turn is not used?? (2)





Algorithm for Process P_i (P_j)

P_i

```
while (true){  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /* remainder section */  
}
```

P_j

```
while (true){  
    flag[j] = true;  
    turn != i;  
    while (flag[i] && turn == i)  
        ;  
  
    /* critical section */  
  
    flag[j] = false;  
  
    /* remainder section */  
}
```

after P_j 's executing section, $\boxed{\text{turn} = j}$
" P_i cannot execute its critical execution

- * What happens if $\text{flag}[]$ is not used?? (1) \Rightarrow progress requirement is not satisfied.
- * What happens if turn is not used?? (2)





Algorithm for Process P_i (P_j)

P_i

```
while (true) {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;
    /* critical section */

    flag[i] = false;

    /* remainder section */
}
```

P_j

```
while (true) {
    flag[j] = true;
    turn = i;
    while (flag[i] && turn == i)
        ;
    /* critical section */

    flag[j] = false;

    /* remainder section */
}
```

The code shows two processes, P_i and P_j , each using a shared flag array and a turn variable. Both processes enter a critical section by setting their respective flag to true and then checking if the other's flag is also true and if the turn variable matches. If so, it enters an infinite loop. Handwritten annotations highlight these loops with green boxes and the text "infinite loop". The code is enclosed in brackets with matching braces at the end of each process.

- * What happens if flag[] is not used?? (1)
- * What happens if turn is not used?? (2)

bounded waiting &
Progress requirement are
not guaranteed.





Peterson's Solution (Cont.)

Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

P_i enters CS only if:

either `flag[j] = false` or `turn = i`

2. Progress requirement is satisfied

3. Bounded-waiting requirement is met





Peterson's Solution (Cont'd)

Although useful for demonstrating an algorithm, Peterson's Solution is not guaranteed to work on modern architectures.

Understanding why it will not work is also useful for better understanding race conditions.

To improve performance, processors and/or compilers may reorder operations that have no dependencies.

For single-threaded this is ok as the result will always be the same.

For multithreaded the reordering may produce inconsistent or unexpected results!





Peterson's Solution (Cont'd)

Two threads share the data:

```
boolean flag = false;  
int x = 0;
```

Thread 1 performs

```
while (!flag);    ① ①  
print x          ② ③
```

Thread 2 performs

```
x = 100;          ② ④  
flag = true       ③ ⑤  
                  100 ⑥
```

What is the expected output?





Peterson's Solution (Cont'd)

100 is the expected output.

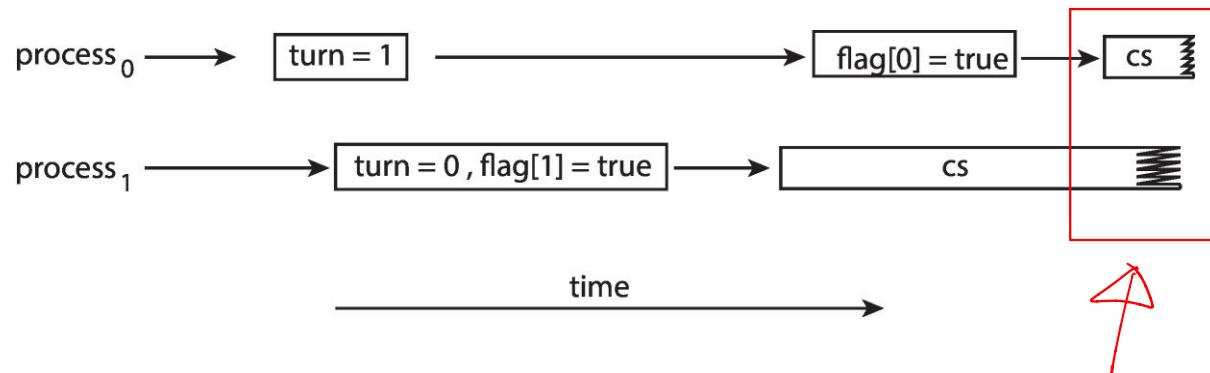
However, the operations for Thread 2 may be reordered:

```
flag = true;  
x = 100;
```

If this occurs, the output may be 0!

상대 차례 & 상대 순서만 있으면 => 양립

The effects of instruction reordering in Peterson's Solution



This allows both processes to be in their critical section at the same time!





Algorithm for Process P_i (P_j)

P_i

```
while (true) {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;
    /* critical section */

    flag[i] = false;

    /* remainder section */
}
```

P_j

```
while (true) {
    flag[j] = true;
    turn = i;
    while (flag[i] && turn == i)
        ;
    /* critical section */

    flag[j] = false;

    /* remainder section */
}
```

reorder

- * What happens if flag[] is not used?? (1)
- * What happens if turn is not used?? (2)





Synchronization Hardware

Many systems provide hardware support for implementing the critical section code.

Uniprocessors – could disable interrupts make non-preemptive

Currently running code would execute without preemption

Generally too inefficient on multiprocessor systems

- ▶ Operating systems using this not broadly scalable

We will look at three forms of hardware support:

1. Memory barriers
2. Hardware instructions
3. Atomic variables





Memory Barriers

Memory models are the memory guarantees a computer architecture makes to application programs.

Memory models may be either:

↗ Can be guaranteed using memory barrier

- **Strongly ordered** – where a **memory modification of one processor is immediately visible to all other processors.** ⇒ reordering instruction is difficult.
- **Weakly ordered** – where a **memory modification of one processor may not be immediately visible to all other processors.**

A **memory barrier** is an instruction that forces any change in memory to be **propagated (made visible) to all other processors.**

⇒ difficult to reorder?





Memory Barrier

We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100:

Thread 1 now performs

```
① while (!flag)
  ↓   memory_barrier();
② print x
```

Thread 2 now performs

```
① x = 100;
  ↓ memory_barrier();
② flag = true
```





Hardware Instructions

Special hardware instructions that allow us to either *test-and-modify* the content of a word, or two *swap* the contents of two words atomically (uninterruptibly.)

Test-and-Set instruction

Compare-and-Swap instruction





test_and_set Instruction

in HW

not interruptable

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = true;
    return rv;
}
```



Run in a
atomic way

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to **true**





Solution using test_and_set()

Shared boolean variable **lock**, initialized to **false**

Solution:

```
do {
```

```
    while (test_and_set(&lock))
```

```
        ; /* do nothing */
```

return true \Rightarrow other process already has lock \Rightarrow wait for lock release

return false \Rightarrow no process has lock yet. \Rightarrow get lock &

entry section

enter critical section

```
    /* critical section */
```

```
lock = false; exit section
```

```
    /* remainder section */
```

```
} while (true);
```





compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```



1. Executed atomically
2. Returns the original value of passed parameter **value**
3. Set the variable **value** the value of the passed parameter **new_value** but only if ***value == expected** is true. That is, the swap takes place only under this condition.





Solution using compare_and_swap

Shared integer `lock` initialized to 0;

Solution:

`while (true) {` return 1 \Rightarrow other process already has lock \Rightarrow wait for lock release

`while (compare_and_swap(&lock, 0, 1) != 0)` return 0 \Rightarrow no process has lock yet \Rightarrow get lock & enter critical section

`; /* do nothing */` entry section

`/* critical section */`

`lock = 0;` exit section

`/* remainder section */`

`}`





Bounded-waiting Mutual Exclusion with compare-and-swap

```
while (true) {
    waiting[i] = true;
    key = 1;
    (A) ←———— while (waiting[i] && key == 1)
        key = compare_and_swap(&lock, 0, 1);
        waiting[i] = false;           entry
    /* critical section */
    j = (i + 1) % n;             exit
    while ((j != i) && !waiting[j]) Pj
        j = (j + 1) % n;          ↗ Circular queue
    if (j == i)                  // Finding a process waiting
        lock = 0;
    else
        waiting[j] = false;      // Opening lock if no process is waiting
    /* remainder section */
}
```

// Setting process out of the while loop (A)

a thread can get its turn after waiting for at most $n-1$ threads





Atomic Variables

Typically, instructions such as compare-and-swap are used as building blocks for other synchronization tools.

One tool is an **atomic variable** that provides *atomic* (uninterruptible) updates on basic data types such as integers and booleans.

For example, the `increment()` operation on the atomic variable `sequence` ensures `sequence` is incremented without interruption:

```
increment(&sequence);
```





Atomic Variables (Cont'd)

The `increment()` function can be implemented as follows:

```
void increment	atomic_int *v)
{
    int temp;
    do {
        temp = *v;
        while (temp != (compare_and_swap(v, temp, temp+1)))
    }
}
```

Annotations on the code:

- A red curly brace on the left side of the code is labeled "SW implementation".
- An annotation above the variable `temp` says "shared variable".
- An annotation above the `while` loop says "HW implementation \Rightarrow atomic".
- A green curly brace on the right side of the `while` loop is labeled "check change of value".
- A green arrow points from the text "other process might change the value of temp before while statement" to the `while` condition.

Other process might change the value of temp before while statement





Mutex Locks

mutual exclusive

Previous solutions are complicated and generally inaccessible to application programmers

OS designers build software tools to solve critical section problem

Simplest is mutex lock

Protect a critical section by first **acquire()** a lock then **release()** the lock

Boolean variable indicating if lock is available or not

Calls to **acquire()** and **release()** must be atomic

Usually implemented via hardware atomic instructions such as compare-and-swap.

But this solution requires **busy waiting**

This lock therefore called a **spinlock**





Solution to Critical-section Problem Using Locks

```
while (true) {
    acquire lock
    critical section
    release lock
    remainder section
}
```





Mutex Lock Definitions

```
initial value of available = true.  
↳ availability of lock  
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false; ;  
}  
  
release() {  
    available = true;  
}
```

These two functions must be implemented atomically.
Both test-and-set and compare-and-swap can be used to implement these functions.





Semaphore

Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.

Semaphore **S** – integer variable

★ Can only be accessed via two indivisible (atomic) operations

wait() and **signal()**

► (Originally called **P()** and **V()**)

Definition of the **wait()** operation

```
wait(S) { request resource, entry section
          while (S <= 0)
              ; // busy wait ⇒ unnecessary usage of CPU
          S--;
      }
```

Definition of the **signal()** operation

```
signal(S) { release resource, exit section
            S++;
        }
```





Semaphore Usage

Counting semaphore – integer value can range over an unrestricted domain

Binary semaphore – integer value can range only between 0 and 1

Same as a **mutex lock**

Can solve various synchronization problems

Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “synch” initialized to 0

P_1 :

① $s_1;$
② signal(synch); *synch = 1*

P_2 :

① wait(synch); *out of loop*
② $s_2;$

Can implement a counting semaphore S as a binary semaphore





Semaphore Implementation

Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time

Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section

Could now have **busy waiting** in critical section implementation

- ▶ But implementation code is short
- ▶ Little busy waiting if critical section rarely occupied

Note that applications may spend lots of time in critical sections and therefore this is not a good solution





Semaphore Implementation with no Busy waiting

With each semaphore there is an associated **waiting queue**

Each entry in a waiting queue has two data items:

- value (of type integer)

- pointer to next record in the list

Two operations:

- block** – place the process invoking the operation on the appropriate waiting queue

- wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct {
    int value; ⇒ status of resource
    struct process *list; ⇒ wait queue
} semaphore;
```

(+) available
(○) full usage
(⊖) waiting processes





Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) { entry
    S->value--;
    if (S->value < 0) { waiting processes exist
        add this process to S->list; put into waiting queue.
        block(); block current process
    }
}

signal(semaphore *S) { exit
    S->value++;
    if (S->value <= 0) { // at least one process is waiting
        remove a process P from S->list;
        wakeup(P); get out of block()
    }
}
```

Handwritten annotations:

- entry: written above the first brace.
- waiting processes exist: handwritten note next to the condition in the if block.
- put into waiting queue.: handwritten note next to the list modification in the if block.
- block current process: handwritten note next to the block call in the if block.
- exit: written above the second brace.
- get out of block(): handwritten note next to the wakeup call in the if block.
- & waiting queue: handwritten note next to the brace closing the if block.





Problems with Semaphores

Incorrect use of semaphore operations: *most of processes can be blocked.*

`wait (mutex) ... wait (mutex)`

Omitting of `wait (mutex)` and/or `signal (mutex)`

These – and others – are examples of what can occur when semaphores and other synchronization tools are used incorrectly.





Monitors

A high-level abstraction that provides a convenient and effective mechanism for process synchronization

Abstract data type, internal variables only accessible by code within the procedure

Only one process may be active within the monitor at a time

Pseudocode syntax of a monitor:

```
monitor monitor-name
{
    // shared variable declarations
    function P1 (...) { .... }

    function P2 (...) { .... }

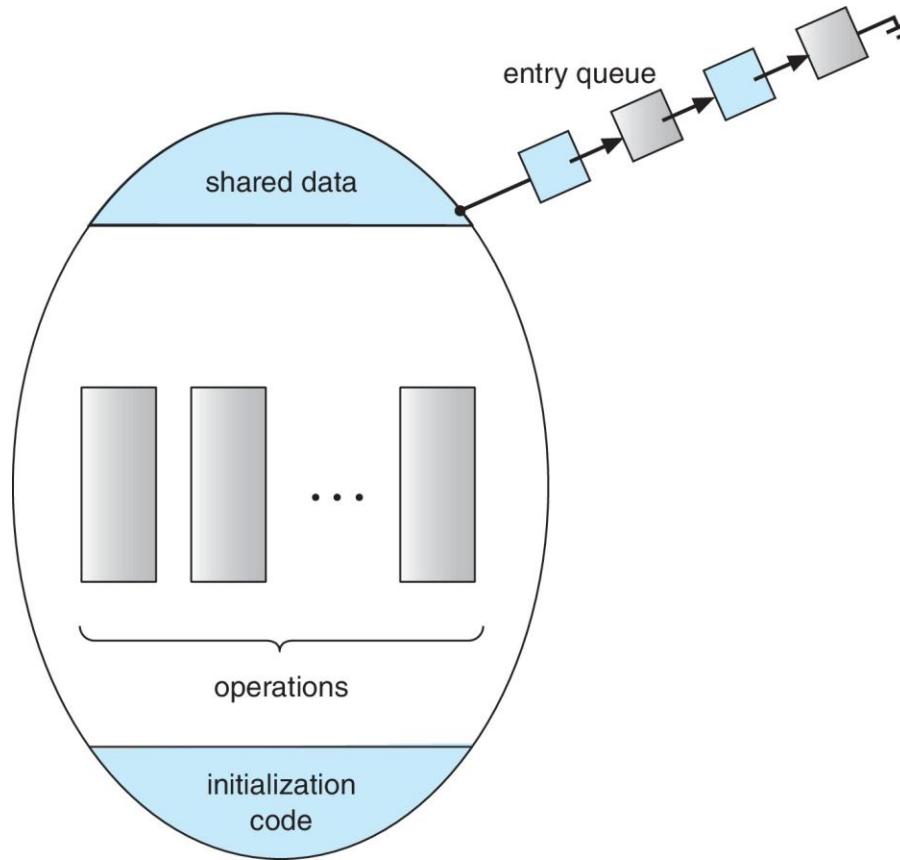
    function Pn (...) { ..... }

    initialization code (...) { ... }
}
```





Schematic view of a Monitor





Condition Variables

```
condition x, y;
```

Two operations are allowed on a condition variable:

x.wait() – a process that invokes the operation is suspended until **x.signal()**

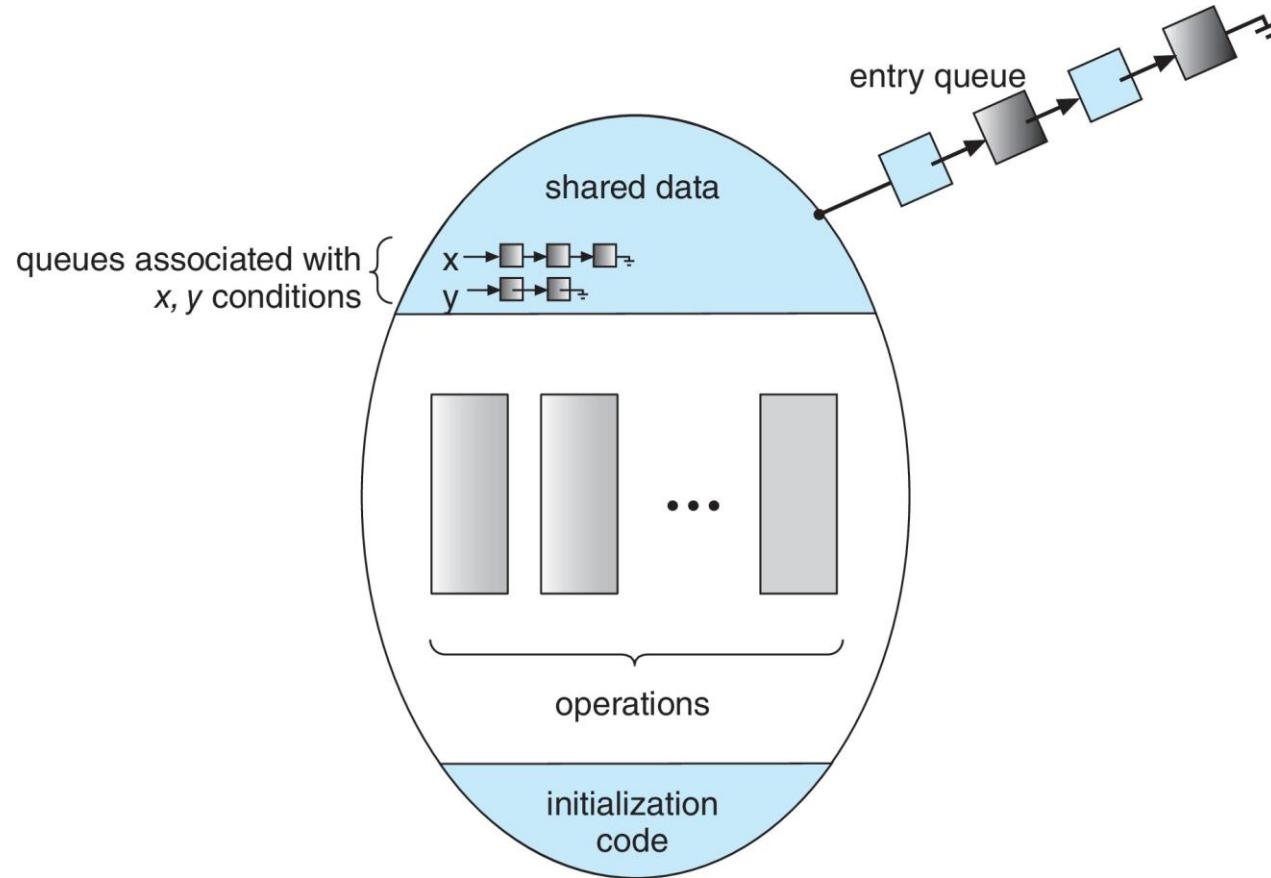
x.signal() – resumes one of processes (if any) that invoked **x.wait()**

- ▶ If no **x.wait()** on the variable, then it has no effect on the variable





Monitor with Condition Variables





Condition Variables Choices

If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?

Both Q and P cannot execute in parallel. If Q is resumed, then P must wait

Options include

Signal and wait – P waits until Q either leaves the monitor or it waits for another condition

Signal and continue – Q waits until P either leaves the monitor or it waits for another condition

language implementer can decide

Monitors implemented in Concurrent Pascal compromise

► P executing signal immediately leaves the monitor, Q is resumed

Implemented in other languages including Mesa, C#, Java





Monitor Implementation Using Semaphores

Variables

```
semaphore mutex; // (initially = 1)  
semaphore next; // (initially = 0)  
int next_count = 0;
```

*first process pass
basically blocked
after signal()*

Each function *F* will be replaced by

```
wait(mutex);  
...  
body of F;  
...  
if (next_count > 0)  
    signal(next)  
else  
    signal(mutex);
```

Mutual exclusion within a monitor is ensured





Monitor Implementation – Condition Variables

For each condition variable **x**, we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

The operation **x.wait()** can be implemented as:

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```





Monitor Implementation (Cont.)

The operation `x.signal()` can be implemented as:

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```





Resuming Processes within a Monitor

If several processes queued on condition variable `x`, and
`x.signal()` is executed, which process should be resumed?

FCFS frequently not adequate

conditional-wait construct of the form `x.wait(c)`

Where `c` is priority number

Process with lowest number (highest priority) is scheduled next





Single Resource allocation

Allocate a single resource among competing processes using priority numbers that specify the maximum time a process plans to use the resource

smallest time, highest priority

```
R.acquire(t);  
...  
access the resource;  
...  
  
R.release;
```

Where R is an instance of type **ResourceAllocator**





A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = true;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = false;
    }
}
```





Liveness

Processes may have to wait indefinitely while trying to acquire a synchronization tool such as a mutex lock or semaphore.

Waiting indefinitely violates the progress and bounded-waiting criteria discussed at the beginning of this chapter.

Liveness refers to a set of properties that a system must satisfy to ensure processes make progress.

Indefinite waiting is an example of a liveness failure.

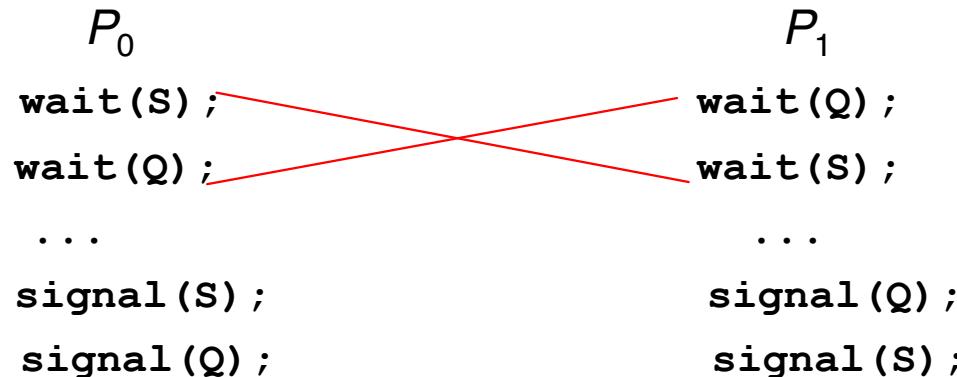




Liveness (Cont'd)

Deadlock – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let S and Q be two semaphores initialized to 1



Consider if P_0 executes $\text{wait}(S)$ and P_1 $\text{wait}(Q)$. When P_0 executes $\text{wait}(Q)$, it must wait until P_1 executes $\text{signal}(Q)$

However, P_1 is waiting until P_0 execute $\text{signal}(S)$.

Since these $\text{signal}()$ operations will never be executed, P_0 and P_1 are **deadlocked**.





Liveness

Other forms of deadlock:

Starvation – indefinite blocking

priority based scheduling
low priority cannot get resource.

- ▶ A process may never be removed from the semaphore queue in which it is suspended

Priority Inversion – Scheduling problem when lower-priority process holds a lock needed by higher-priority process

- ▶ Solved via **priority-inheritance protocol**





Priority Inheritance Protocol

Consider the scenario with three processes **P1**, **P2**, and **P3**. **P1** has the highest priority, **P2** the next highest, and **P3** the lowest. Assume **P3** is assigned a resource **R** that **P1** wants. Thus, **P1** must wait for **P3** to finish using the resource. However, **P2** becomes runnable and preempts **P3**. What has happened is that **P2** - a process with a lower priority than **P1** - has indirectly prevented **P3** from gaining access to the resource.

To prevent this from occurring, a **priority inheritance protocol** is used. This simply allows the priority of the highest thread waiting to access a shared resource to be assigned to the thread currently using the resource. Thus, the current owner of the resource is assigned the priority of the highest priority thread wishing to acquire the resource.



End of Chapter 6

