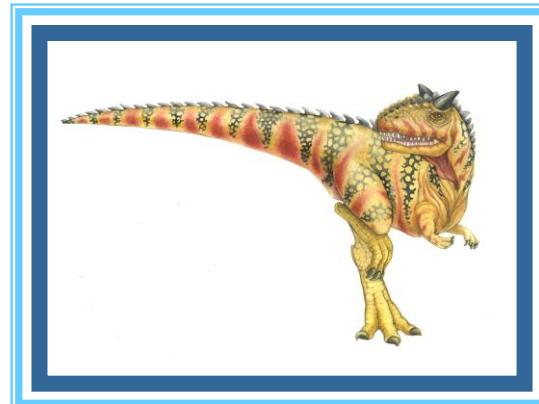


Chapter 8: Deadlocks





Chapter 8: Deadlocks

System Model

Deadlock in Multithreaded Applications

Deadlock Characterization

Methods for Handling Deadlocks

 Deadlock Prevention

 Deadlock Avoidance

 Deadlock Detection

 Recovery from Deadlock





Chapter Objectives

- Illustrate how deadlock can occur when mutex locks are used
- Define the four necessary conditions that characterize deadlock
- Identify a deadlock situation in a resource allocation graph
- Evaluate the four different approaches for preventing deadlocks
- Apply the banker's algorithm for deadlock avoidance
- Apply the deadlock detection algorithm
- Evaluate approaches for recovering from deadlock





System Model

System consists of resources

Resource types R_1, R_2, \dots, R_m

CPU cycles, memory space, I/O devices

Each resource type R_i has W_i instances.

Each process utilizes a resource as follows:

1 request

2 use

3 release





Deadlock in Multithreaded Application

```
pthread_mutex_t first_mutex;  
pthread_mutex_t second_mutex;  
  
pthread_mutex_init(&first_mutex,NULL);  
pthread_mutex_init(&second_mutex,NULL);
```





Deadlock in Multithreaded Application

```
1st thread
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex); 1
    pthread_mutex_lock(&second_mutex); 2
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex); 2
    pthread_mutex_unlock(&first_mutex); 1

    pthread_exit(0);
}

2nd thread
/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex); 2
    pthread_mutex_lock(&first_mutex); 1
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex); 1
    pthread_mutex_unlock(&second_mutex); 2

    pthread_exit(0);
}
```

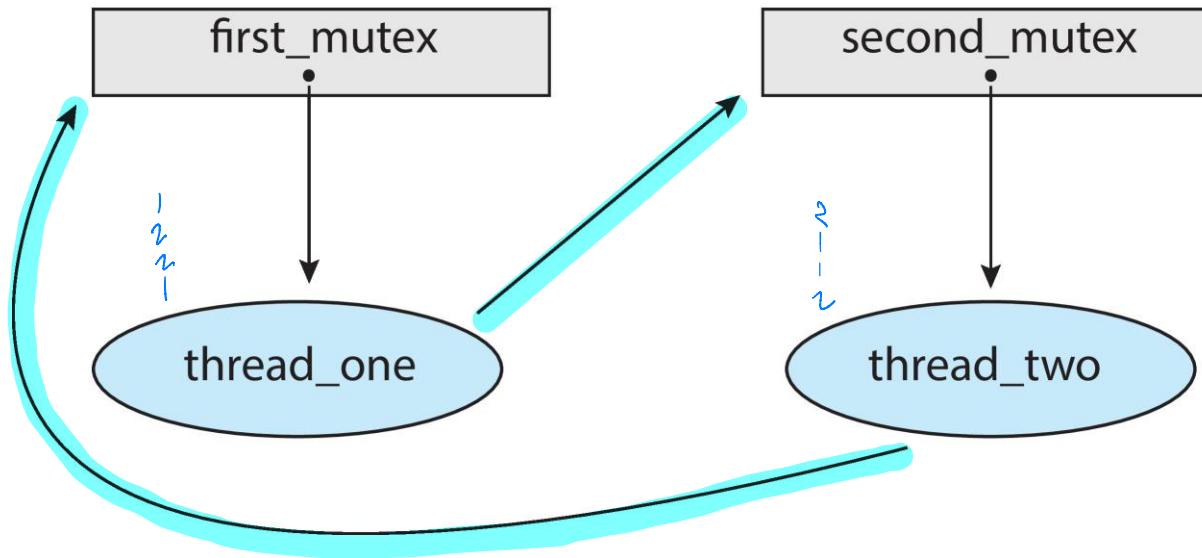




Deadlock in Multithreaded Application

Deadlock is possible if thread 1 acquires `first_mutex` and thread 2 acquires `second_mutex`. Thread 1 then waits for `second_mutex` and thread 2 waits for `first_mutex`. \Rightarrow deadlock

Can be illustrated with a **resource allocation graph**:





Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

Mutual exclusion: only one process at a time can use a resource

Hold and wait: a process holding at least one resource is waiting to acquire additional resources held by other processes

No preemption: a resource can be released only voluntarily by the process holding it, after that process has completed its task

Circular wait: there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .





Resource-Allocation Graph

A set of vertices V and a set of edges E .

V is partitioned into two types:

$P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system

$R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system

request edge – directed edge $P_i \rightarrow R_j$

assignment edge – directed edge $R_j \rightarrow P_i$





Resource Allocation Graph Example

One instance of R1

Two instances of R2

One instance of R3

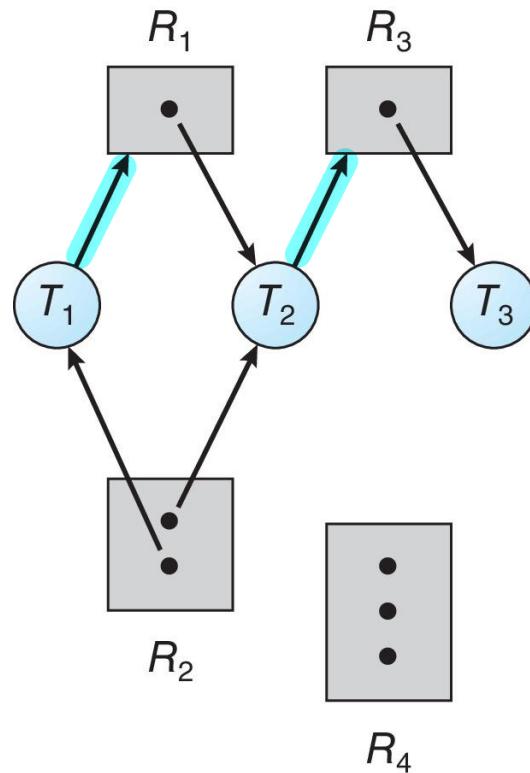
Three instances of R4

T1 holds one instance of R2 and is waiting for an instance of R1

T2 holds one instance of R1, one instance of R2, and is waiting for an instance of R3

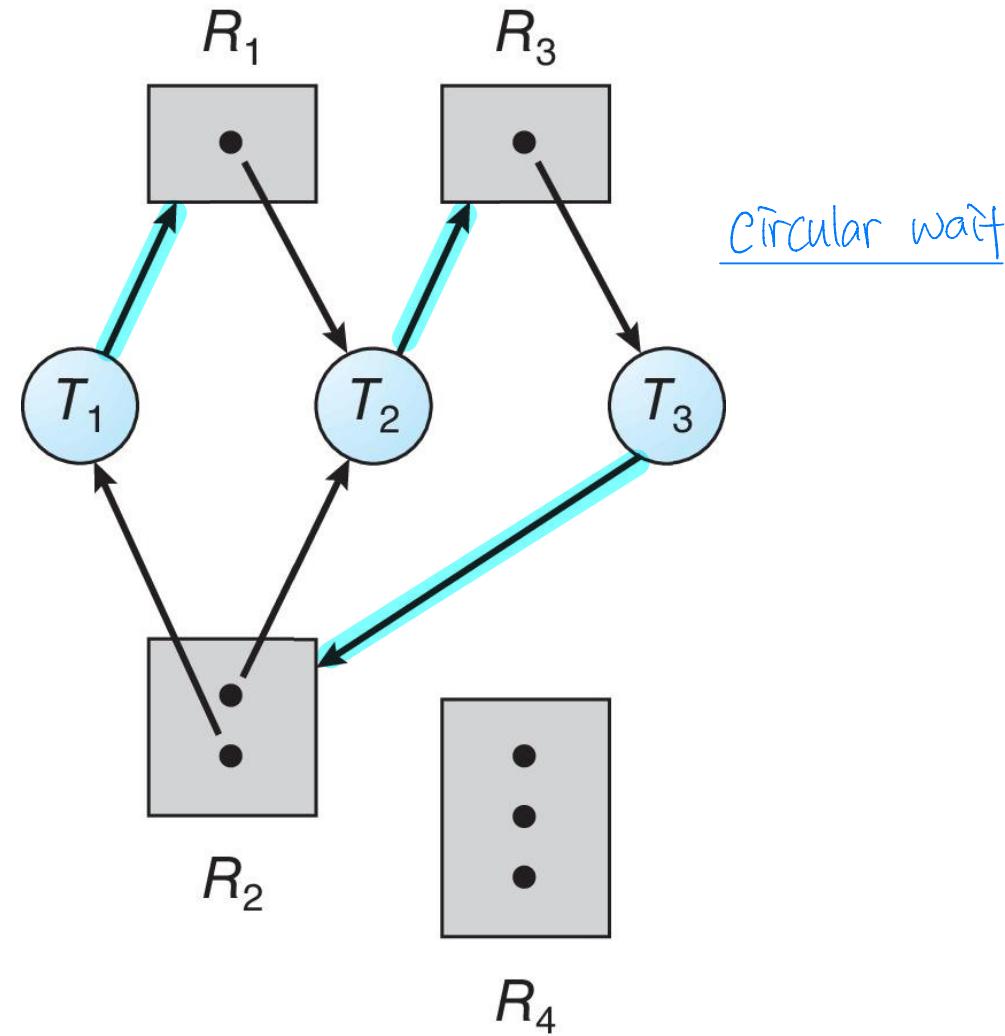
T3 is holding one instance of R3

No Deadlock





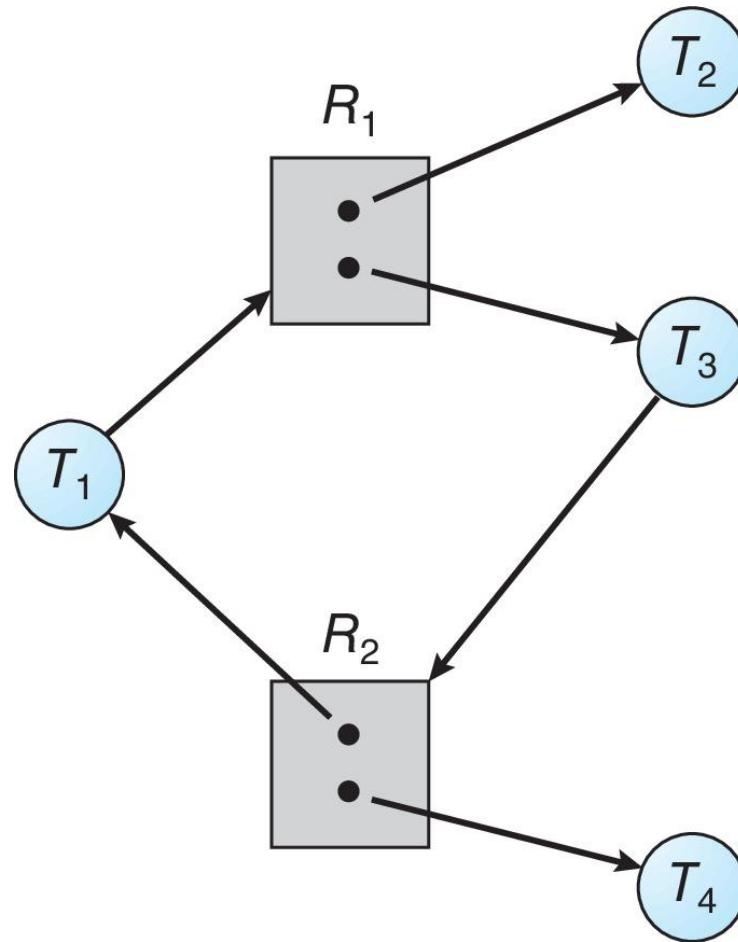
Resource Allocation Graph With A Deadlock





Graph With A Cycle But No Deadlock

$T_4 \rightarrow T_3 \rightarrow T_1$
 $T_2 \rightarrow T_1 \rightarrow T_3$





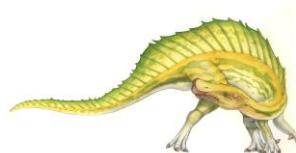
Basic Facts

If graph contains no cycles \Rightarrow no deadlock

If graph contains a cycle \Rightarrow

if only one instance per resource type, then deadlock

if several instances per resource type, possibility of deadlock





Methods for Handling Deadlocks

Ensure that the system will **never** enter a deadlock state:

- Deadlock prevention
- Deadlock avoidance

Allow the system to enter a deadlock state and then recover

- Ignore the problem and pretend that deadlocks never occur in the system.





Deadlock Prevention

⇒ Severe performance degradation.

Invalidate one of the four necessary conditions for deadlock:

Mutual Exclusion – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources

Hold and Wait – must guarantee that whenever a process requests a resource, it does not hold any other resources

Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.

Low resource utilization; starvation possible





Deadlock Prevention (Cont.)

No Preemption –

If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released

Preempted resources are added to the list of resources for which the process is waiting

Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration





Circular Wait

Invalidating the circular wait condition is most common.

Simply assign each resource (i.e. mutex locks) a unique number.

Resources must be acquired in order.

If:

```
first_mutex = 1  
second_mutex = 5
```

code for **thread_two** could not be
written as follows:

MUST Increasing order

```
/* thread_one runs in this function */  
void *do_work_one(void *param)  
{  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
  
    pthread_exit(0);  
}  
  
/* thread_two runs in this function */  
void *do_work_two(void *param)  
{  
    pthread_mutex_lock(&second_mutex); 5  
    pthread_mutex_lock(&first_mutex); 1  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&first_mutex);  
    pthread_mutex_unlock(&second_mutex);  
  
    pthread_exit(0);  
}
```





Deadlock Avoidance

maintain overall system performance

Requires that the system has some additional *a priori* information available
given in advance

Simplest and most useful model requires that each process declare the **maximum number** of resources of each type that it may need

The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition

Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes





Safe State

When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state

System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$

That is:

If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished

When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate

When P_i terminates, P_{i+1} can obtain its needed resources, and so on





Basic Facts

If a system is **in safe state** \Rightarrow **no deadlocks**

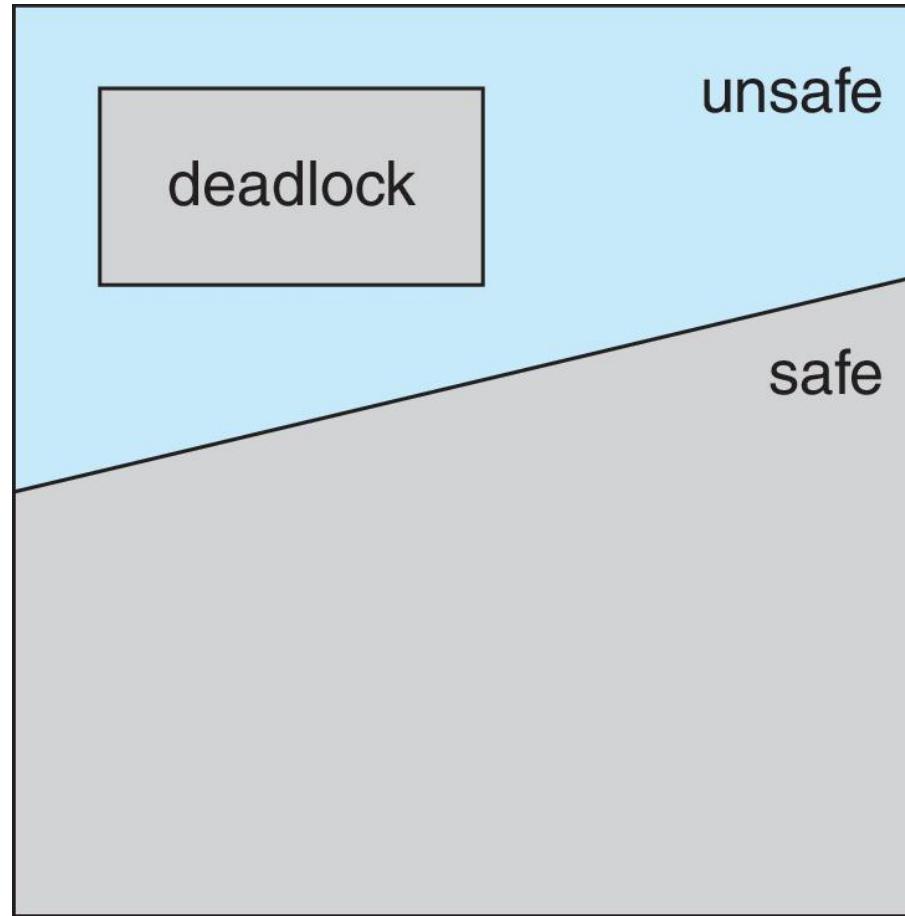
If a system is **in unsafe state** \Rightarrow **possibility of deadlock**

- **Avoidance** \Rightarrow ensure that a system will **never enter an unsafe state**.





Safe, Unsafe, Deadlock State





Avoidance Algorithms

- Single instance of a resource type

- Use a resource-allocation graph

- Multiple instances of a resource type

- Use the Banker's Algorithm





Single instance per resource

Resource-Allocation Graph Scheme

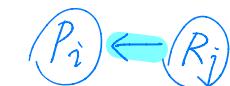
Claim edge $P_i \rightarrow R_j$ indicated that process P_i may request resource R_j ; represented by a dashed line



Claim edge converts to **request edge** when a process requests a resource (solid line)



Request edge converted to an **assignment edge** when the resource is allocated to the process



When a resource is released by a process, assignment edge reconverts to a claim edge



Resources must be claimed *a priori* in the system





Resource-Allocation Graph Algorithm

Suppose that process P_i requests a resource R_j

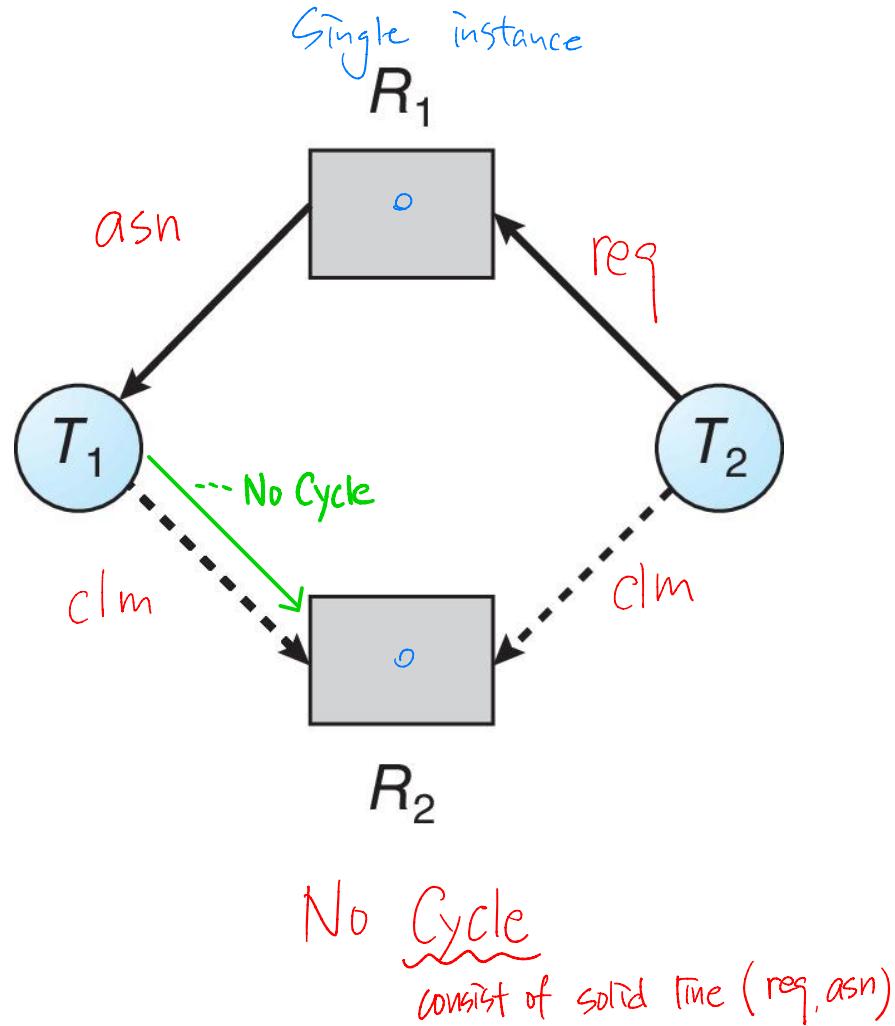
The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

$$\begin{array}{c} P_i \rightarrow R_j \\ \Downarrow \text{No Cycle!} \\ P_i \leftarrow R_j \end{array}$$



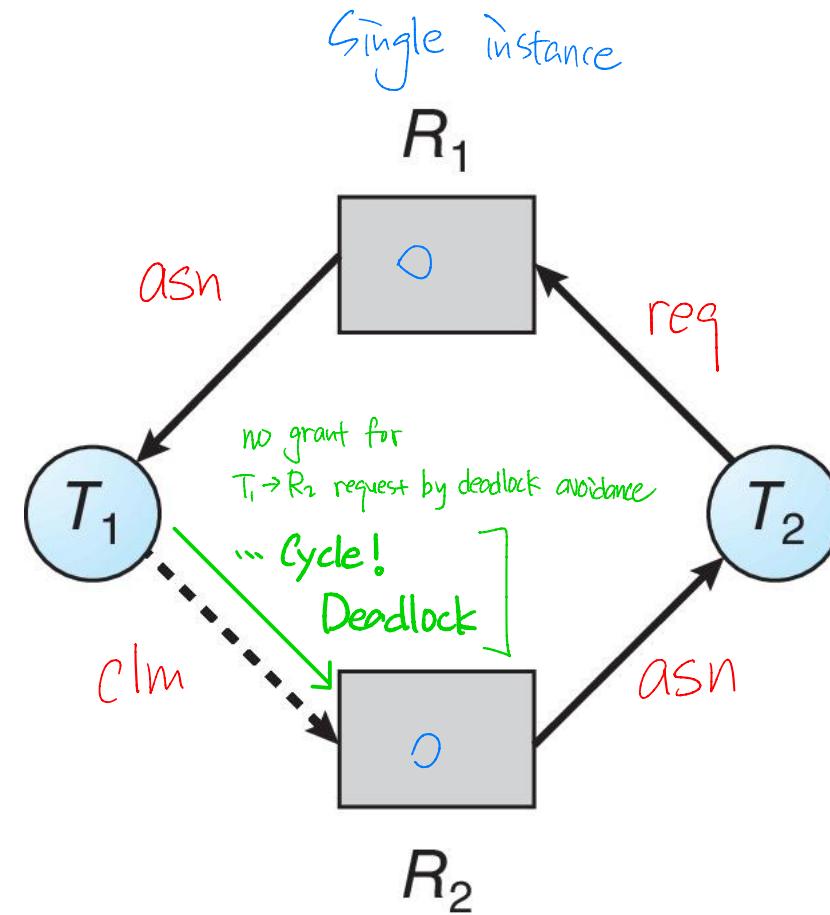


Resource-Allocation Graph





Unsafe State In Resource-Allocation Graph



No Cycle





Banker's Algorithm

Multiple instances of resources

Each process must *a priori* claim maximum use

When a process requests a resource it may have to wait

When a process gets all its resources it must return them in a finite amount of time
eventually





Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- 1 **Available:** Vector of length m . If $\text{available}[j] = k$, there are k instances of resource type R_j available
- 2 **Max:** $n \times m$ matrix. If $\text{Max}[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- 3 **Allocation:** $n \times m$ matrix. If $\text{Allocation}[i,j] = k$ then P_i is currently allocated k instances of R_j
- 4 **Need:** $n \times m$ matrix. If $\text{Need}[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

Worst scenario : process requests max instance
⇒ still safe ?





Safety Algorithm

1. Let **Work** and **Finish** be vectors of length m and n , respectively.
Initialize:

Work = Available

Finish [i] = false for $i = 0, 1, \dots, n-1$

2. Find an i such that both:

(a) **Finish [i] = false** \uparrow not finished process P_i

(b) **Need_i ≤ Work** \rightarrow we can allocate $need_i$ resource right now.

If no such i exists, go to step 4

3. **Work = Work + Allocation_i**

Finish[i] = true

go to step 2

the resources are released after termination.

P_i execution finished,

4. If **Finish [i] == true** for all i , then the system is in a safe state

else, in unsafe state





Resource-Request Algorithm for Process P_i

Request_i = request vector for process P_i . If $\text{Request}_i[j] = k$ then process P_i wants k instances of resource type R_j

1. If $\text{Request}_i \leq \text{Need}_i$, go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If $\text{Request}_i \leq \text{Available}$, go to step 3. Otherwise P_i must wait,
and holding resource since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$\text{Available} = \text{Available} - \text{Request}_i;$$

$$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i;$$

$$\text{Need}_i = \text{Need}_i - \text{Request}_i;$$

If safe \Rightarrow the resources are allocated to P_i

If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

by safety algorithm





Example of Banker's Algorithm

5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>	<u>Available</u>	<u>Need = Max - Allocation</u>
	A	B	C	A	B	C
P_0	0 1 0			7 5 3	3 3 2	7 4 3
P_1	2 0 0			3 2 2		1 2 2
P_2	3 0 2			9 0 2		6 0 0
P_3	2 1 1			2 2 2		0 1 1
P_4	0 0 2			4 3 3		4 3 1





Example (Cont.)

The content of the matrix **Need** is defined to be **Max – Allocation**

<u>Allocation</u>			<u>Need</u>			<u>Available</u>				
	A	B	C	A	B	C	A	B	C	
P_0	0	1	0	P_0	7	4	3	?	?	
P_1	2	0	0	P_1	1	2	2	allocate [Need], release [Need] & [Allocation]		
P_2	3	0	2	P_2	6	0	0	$(3, 3, 2) \rightarrow (2, 1, 0) \rightarrow (5, 3, 2)$		
P_3	2	1	1	P_3	0	1	1			
P_4	0	0	2	P_4	4	3	1			

The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria *not unique sol*

Available: $(3,3,2) \rightarrow (5,3,2) \rightarrow (7,4,3) \rightarrow (7,4,5) \rightarrow (10,4,7) \rightarrow (10,5,7)$





Example: P_1 Request (1,0,2)

Check that $\text{Request} \leq \text{Available}$ (that is, $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$) *granted*

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>			initial = $(3,3,2) - (1,0,2)$ $= (2,3,0)$
	A	B	C	A	B	C	A	B	C	
P_0	0	1	0	7	4	3	7	5	5	⊕
P_1	3	0	2	0	2	0	5	3	2	①
P_2	3	0	2	6	0	0	10	5	7	②
P_3	2	1	1	0	1	1	7	4	3	③
P_4	0	0	2	4	3	1	7	4	5	④

Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement

Can request for (3,3,0) by P_4 be granted?

Can request for (0,2,0) by P_0 be granted?





Example: P_1 Request (1,0,2)

additional

Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$)

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>			
	A	B	C	A	B	C	A	B	C	initial
P_0	0	3	0	7	2	3	2	1	0	$(3,3,2) \rightarrow (2,3,0) \rightarrow (2,1,0)$
P_1	3	0	2	0	2	0	5	3	2	
P_2	3	0	2	6	0	0	2	0	0	
P_3	2	1	1	0	1	1	7	4	3	
P_4	0	0	2	4	3	1	7	4	5	

Executing safety algorithm shows that sequence $< P_1, P_3, P_4, P_0, P_2 >$ satisfies safety requirement

Can request for (3,3,0) by P_4 be granted? No. $(2,3,0) \prec (3,3,0)$

Can request for (0,2,0) by P_0 be granted? No. $(2,3,0) \succ (0,2,0) \Rightarrow (2,1,0)$





Deadlock Detection

- Allow system to enter deadlock state
 - Detection algorithm
 - Recovery scheme





Single Instance of Each Resource Type

Maintain **wait-for** graph

Nodes are processes

$P_i \rightarrow P_j$ if P_i is waiting for P_j

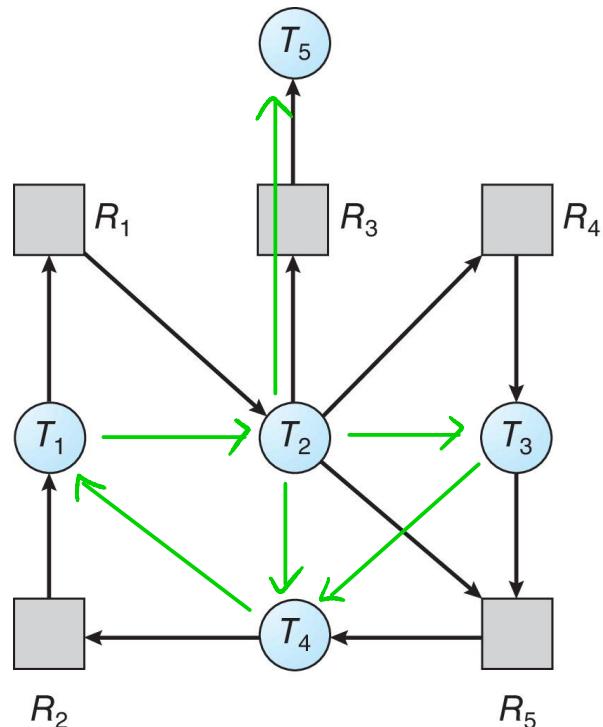
Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock \Rightarrow then recovery!

An algorithm to detect a cycle in a graph requires an order of n^2 $O(n^2)$ operations, where n is the number of vertices in the graph



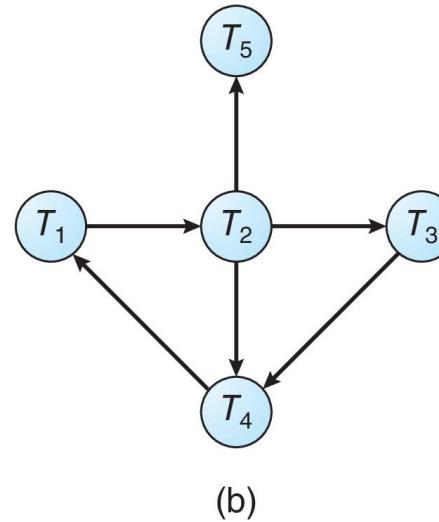


Resource-Allocation Graph and Wait-for Graph



(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph





Several Instances of a Resource Type

Available: A vector of length m indicates the number of available resources of each type

Allocation: An $n \times m$ matrix defines the number of resources of each type currently allocated to each process

Request: An $n \times m$ matrix indicates the current request of each process. If $\text{Request}[i][j] = k$, then process P_i is requesting k more instances of resource type R_j .

No "Max" \Rightarrow no ability to decide grant of request like deadlock avoidance





Detection Algorithm

1. Let **Work** and **Finish** be vectors of length m and n , respectively
Initialize:

(a) **Work = Available** $O(m)$

(b) For $i = 1, 2, \dots, n$, if **Allocation_i ≠ 0**, then
Finish[i] = false; otherwise, **Finish[i] = true** $O(n)$

2. Find an index i such that both: $O(n)$

(a) **Finish[i] == false**

(b) **Request_i ≤ Work** $O(m)$

]

$O(n \times m)$

If no such i exists, go to step 4





Detection Algorithm (Cont.)

3. $\text{Work} = \text{Work} + \text{Allocation}_i$, $O(n)$
 $\text{Finish}[i] = \text{true}$, $O(1)$] $O(n)$
go to step 2
4. If $\text{Finish}[i] == \text{false}$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $\text{Finish}[i] == \text{false}$, then P_i is deadlocked
 $O(n)$

Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state

$$\begin{aligned}& \boxed{\text{Step 2}} \times \boxed{\text{Step 3}} \\& \Rightarrow O(n \times m) \times O(n) = O(m \cdot n^2)\end{aligned}$$





Example of Detection Algorithm

Five processes P_0 through P_4 ; three resource types A (7 instances), B (2 instances), and C (6 instances)

Snapshot at time T_0 :

		<u>Allocation</u>	<u>Request</u>	<u>Available</u>	
		A B C	A B C	A B C	
1	P_0	0 1 0	0 0 0	0 0 0	0 1 0 ①
4	P_1	2 0 0	2 0 2 ↙	7 2 4	④
2	P_2	3 0 3	0 0 0 ↙	3 1 3	②
3	P_3	2 1 1	1 0 0 ↙	5 2 4	③
5	P_4	0 0 2	0 0 2 ↙	7 2 6	⑤

Sequence $\langle P_0, P_2, \underbrace{P_3, P_1, P_4}_{\text{order is not unique}} \rangle$ will result in $Finish[i] = \text{true}$ for all i





Example (Cont.)

P_2 requests an additional instance of type C

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	P_0	0	0	0	0	0
P_1	2	0	0	P_1	2	0	2	x	x
P_2	3	0	3	P_2	0	0	1	x	x
P_3	2	1	1	P_3	1	0	0	x	x
P_4	0	0	2	P_4	0	0	2	x	x

State of system?

Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes' requests

Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4





Detection-Algorithm Usage

When, and how often, to invoke depends on:

- How often a deadlock is likely to occur?
- How many processes will need to be rolled back?
 - ▶ one for each disjoint cycle

If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.





Recovery from Deadlock: Process Termination

Abort all deadlocked processes

Abort one process at a time until the deadlock cycle is eliminated

In which order should we choose to abort?

1. Priority of the process low \rightarrow abort
2. How long process has computed, and how much longer to completion short \rightarrow abort
3. Resources the process has used many \rightarrow abort
4. Resources process needs to complete many \rightarrow abort
5. How many processes will need to be terminated minimize
6. Is process interactive or batch? batch \rightarrow abort





better

Recovery from Deadlock: Resource Preemption

Selecting a victim – minimize cost

Rollback – return to some safe state, restart process for that state

Starvation – same process may always be picked as victim, include number of rollback in cost factor \Rightarrow aging



End of Chapter 8

